

Batch univariate mixture model

Stephen Coleman

20/01/2021

Batch modelling

We write the basic mixture model for independent items $X = (x_1, \dots, x_N)$ as

$$x_n \sim \sum_{k=1}^K \pi_k f(x_n | \theta_k) \quad \text{independently for } n = 1, \dots, N \quad (1)$$

where $f(\cdot | \theta)$ is some family of densities parametrised by θ . A common choice is the Gaussian density function, with $\theta = (\mu, \sigma^2)$ (as in our simulation study). K , the number of subgroups in the population, $\{\theta_k\}_{k=1}^K$, the component parameters, and $\pi = (\pi_1, \dots, \pi_K)$, the component weights are the objects to be inferred. In the context of *clustering*, such a model arises due to the belief that the population from which the random sample under analysis has been drawn consists of K unknown groups proportional to π . In this setting it is natural to include a latent *allocation variable*, $c = (c_1, \dots, c_N)$, to indicate which group each item is drawn from, with each non-empty component of the mixture corresponds to a cluster. The model is

$$\begin{aligned} p(c_n = k) &= \pi_k \quad \text{for } k = 1, \dots, K, \\ x_n | c_n &\sim f(x_n | \theta_k) \quad \text{independently for } n = 1, \dots, N. \end{aligned} \quad (2)$$

The joint model can then be written

$$p(X, c, K, \pi, \theta) = p(X | c, \pi, K, \theta) p(\theta | c, \pi, K) p(c | \pi, K) p(\pi | K) p(K).$$

We will focus upon the Gaussian kernel with a mean parameter of μ and a precision of τ , so henceforth $\theta = (\mu, \tau)$.

We can extend this model for B batches, introducing an observed *batch variable*, $b = (b_1, \dots, b_N)$, to indicate which batch the n^{th} individual comes from

$$\begin{aligned} p(c_n = k) &= \pi_k \quad \text{for } k = 1, \dots, K, \\ x_n | c_n, b_n &\sim f(x_n | \mu_k + m_b, \tau_k \times t_b) \quad \text{independently for } n = 1, \dots, N. \end{aligned} \quad (3)$$

This model tries to leverage all the data available - the batch coefficients are not cluster specific and thus there is some complex dependencies. We assume prior distributions:

$$\mu_k \sim \mathcal{N}(\xi, \kappa \tau_k) \quad (4)$$

$$\tau_k \sim \text{Ga}(\alpha, \beta) \quad (5)$$

$$m_b \sim \mathcal{N}(\delta, \lambda t_b) \quad (6)$$

$$t_b \sim \text{Ga}(\rho, \theta) \quad (7)$$

and a likelihood function

$$p(X | c, b, \mu, \tau, m, t, \pi) = \prod_{n=1}^N \left(\frac{\tau_{c_n} t_{b_n}}{2\pi} \right)^{\frac{1}{2}} \exp \left\{ -\frac{\tau_{c_n} t_{b_n}}{2} [x_n - (\mu_{c_n} + m_{b_n})]^2 \right\}. \quad (8)$$

Expanding to P independent features we have:

$$p(X|c, b, \mu, \tau, m, t, \pi) = \prod_{n=1}^N \prod_{p=1}^P \left(\frac{\tau_{c_n, p} t_{b_n, p}}{2\pi} \right)^{\frac{1}{2}} \exp \left\{ -\frac{\tau_{c_n, p} t_{b_n, p}}{2} [x_{n, p} - (\mu_{c_n, p} + m_{b_n, p})]^2 \right\}. \quad (9)$$

The joint model is:

$$p(X, c, \mu, \tau, m, t, \pi) = p(\pi|\alpha) \prod_{n=1}^N p(x_n|c_n, b_n, \mu_{c_n}, \sigma_{c_n}^2, a_{b_n}, s_{b_n}^2) p(c_n|\pi) \prod_{k=1}^K p(\mu_k|\mu_0, \lambda_0, \sigma_k^2) p(\sigma_k^2|\gamma, \epsilon) \prod_{b=1}^B p(a_b|\xi, \delta, s_b^2) p(s_b^2|\rho, \theta) \quad (10)$$

$$\Rightarrow p(c, \mu, \tau, m, t|\cdot) \propto \prod_{n=1}^N \left(\frac{\tau_{c_n} t_{b_n}}{2\pi} \right)^{1/2} \exp \left\{ -\frac{\tau_{c_n} t_{b_n}}{2} [x_n - (\mu_{c_n} + m_{b_n})]^2 \right\} \quad (11)$$

$$\times \prod_{k=1}^K \left(\frac{\kappa \tau_k}{2\pi} \right)^{1/2} \exp \left\{ -\frac{\kappa \tau_k}{2} (\mu_k - \xi)^2 \right\} \tau_k^{\alpha-1} \exp \{-\beta \tau_k\} \quad (12)$$

$$\times \prod_{b=1}^B \left(\frac{\lambda t_b}{2\pi} \right)^{1/2} \exp \left\{ -\frac{\lambda t_b}{2} (m_b - \delta)^2 \right\} t_b^{\rho-1} \exp \{-\theta t_b\}. \quad (13)$$

Metropolis-Hastings

The *Metropolis* algorithm (named after the work of Osamu Tezuka) can be used to sample from awkward distributions; i.e. those we only know some kernel to which it is proportional. This algorithm was extended, achronologically, to use asymmetric proposal densities, a result established by the Norman victory at the battle of Hastings and thus the name of the *Metropolis-Hastings* algorithm. We can use these awkward posterior marginal kernels to perform *Metropolis-within-Gibbs*. The relevant kernels are

$$p(t_b|\cdot) \propto t_b^{\frac{1}{2}(N_b+2\rho-1)} \exp \left\{ -\frac{t_b}{2} \left[\sum_{n=1}^N \tau_{c_n} [x_n - (\mu_{c_n} + m_b)]^2 \mathbb{I}(b_n = b) + \lambda(m_b - \delta)^2 + 2\theta \right] \right\}, \quad (14)$$

$$p(\tau_k|\cdot) \propto \tau_k^{\frac{1}{2}(N_k+2\alpha-1)} \exp \left\{ -\frac{\tau_k}{2} \left[\sum_{n=1}^N t_{b_n} [x_n - (\mu_k + m_{b_n})]^2 \mathbb{I}(c_n = k) + \kappa(\mu_k - \xi)^2 + 2\beta \right] \right\}, \quad (15)$$

$$p(m_b|\cdot) \propto \exp \left\{ -\frac{t_b}{2} \left[\sum_{n=1}^N \tau_{c_n} [x_n - (\mu_{c_n} + m_b)]^2 \mathbb{I}(b_n = b) + \lambda(m_b - \delta)^2 \right] \right\}, \quad (16)$$

$$p(\mu_k|\cdot) \propto \exp \left\{ -\frac{\tau_k}{2} \left[\sum_{n=1}^N t_{b_n} [x_n - (\mu_k + m_{b_n})]^2 \mathbb{I}(c_n = k) + \kappa(\mu_k - \xi)^2 \right] \right\}. \quad (17)$$

I use proposal densities of a Gaussian for the mean parameters

$$q(\mu'_k|\mu_k) \sim \mathcal{N}(\mu_k, \sigma_1^2), \quad (18)$$

$$q(m'_b|m_b) \sim \mathcal{N}(m_b, \sigma_1^2). \quad (19)$$

and for the precisions I use one of either a Gamma distribution

$$q(\tau'_k|\tau_k) \sim \text{Ga}(\sigma_2^2 \tau_k, \sigma_2^2), \quad (20)$$

$$q(t'_b|t_b) \sim \text{Ga}(\sigma_2^2 t_b, \sigma_2^2), \quad (21)$$

or a log transform of a Gaussian

$$q(\tau'_k|\tau_k) \sim \log \mathcal{N}(\tau_k, \sigma_2^2), \quad (22)$$

$$q(t'_b|t_b) \sim \log \mathcal{N}(t_b, \sigma_2^2). \quad (23)$$

As the Gaussian distribution is symmetric, this is not present in the acceptance probability, cancelling itself; this does not hold for the precisions. Currently I have that $\sigma_1^2 = 0.5$ and $\sigma_2^2 = 4.0$.

Input:Data X ,The number of iterations, I ,The prior distribution $p(\theta)$,The likelihood function, $p(X|\theta)$,The proposal distribution, $q(\theta)$.**Output:** A vector of accepted values for θ .**begin**

```

    /* initialise theta by drawing from the prior */
     $\theta_0 \sim p(\theta)$ ;
    for  $i = 1$  to  $I$  do
        /* propose a new value */
         $\theta' \sim q(\theta_{i-1})$ ;
        /* calculate the acceptance probability */
         $\alpha = \min(1, \frac{p(X|\theta')p(\theta')q(\theta_{i-1}|\theta')}{p(X|\theta_{i-1})p(\theta_{i-1})q(\theta'|\theta_{i-1})})$ ;
         $u \sim \text{Unif}(0, 1)$ ;
        if  $u < \alpha$  then
            |  $\theta_i \leftarrow \theta'$ ;
        else
            |  $\theta_i \leftarrow \theta_{i-1}$ ;
        end
    end
end

```

Algorithm 1: The Metropolis-Hastings algorithm for Bayesian inference.

R package

I have written an R package to cluster univariate batch data.

```
library(BatchMixtureModel)
```

```
# For some PSM related functions
```

```
library(mdiHelpR)
```

```
##
```

```
## Attaching package: 'mdiHelpR'
```

```
## The following object is masked from 'package:methods':
```

```
##
```

```
##      show
```

```
# For the pipe
```

```
library(magrittr)
```

```
# For data visualisation
```

```
library(ggplot2)
```

```
setMyTheme()
```

I generate data from the model.

```

# Dataset parameters
N <- 200
P <- 1
K <- 2
B <- 5
mean_dist <- 5
batch_dist <- 2
cluster_means <- 1:K * mean_dist
batch_shift <- 1:B * batch_dist
std_dev <- rep(1, K)
batch_var <- rep(1, B)
cluster_weights <- rep(1 / K, K)
batch_weights <- rep(1 / B, B)

# We generate data, acquiring the data with batch effects present and absent
my_data <- generateBatchData(
  N,
  P,
  cluster_means,
  std_dev,
  batch_shift,
  batch_var,
  cluster_weights,
  batch_weights
)

```

We visualise the data.

```


vis_df <- data.frame(
  x = my_data$data,
  x_true = my_data$corrected_data,
  cluster = factor(my_data$cluster_IDs),
  batch = factor(my_data$batch_IDs)
)

# Look at the data with and without batch effects
vis_df %>%
  ggplot(aes(x = x, fill = cluster)) +
  geom_histogram() +
  labs(
    title = "Generated data",
    subtitle = "Batch effects present"
  )

```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Generated data
Batch effects present




```
vis_df %>%  
  ggplot(aes(x = x_true, fill = cluster)) +  
  geom_histogram() +  
  labs(  
    title = "Generated data",  
    subtitle = "Batch effects removed"  
  )
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

Generated data

Batch effects removed



Now let's cluster the data.

```
# Some random initialisation
c_init <- sample(1:K, N, replace = T)

# The proposal windows for the Metropolis steps (currently using log normal for
# precisions)
proposal_window <- 0.5
proposal_window_for_logs <- 0.3

time_0 <- Sys.time()
samples <- sampleMixtureModel(
  matrix(my_data$data, ncol = 1),
  K,
  B,
  c_init = 1,
  my_data$batch_IDs = 1,
  proposal_window,
  proposal_window_for_logs,
  500000,
  1000,
  rep(1, K),
  1
)
time_1 <- Sys.time()
print(time_1 - time_0)
```


```
## Time difference of 1.033424 mins
```

Now let us look at the output.

```
burn <- 1:100

# Make and look at the PSM
psm <- createSimilarityMat(samples$samples[-burn, ]) %>%
  set_rownames(names(my_data$data)) %>%
  set_colnames(names(my_data$data))

annotatedHeatmap(psm, my_data$cluster_IDs,
  col_pal = simColPal(),
  main = "Generated data PSM")
```



```
# A histogram of the cluster weights (one is enough as K=2 and we're on a simplex)
hist(samples$weights[-burn, 1])
```

Histogram of samples\$weights[-burn, 1]

